

# Computer Vision Project 5: TinyYOLO - framework for Handwriting Recognition

Ian Sinclair  
ENCE 3620-1

March 30, 2022  
Dr. Mohammad H. Mahoor

## Abstract

Considered is an implementation of the tinyYOLOv3 (and tiny-YOLOv4) framework using darkflow. YOLO (you only look once) is an advanced image recognition neural network that is designed to recognize classify and draw bounding boxes around multiple objects in an image simultaneously. Here, we re-train the network to detect handwritten digits on a page. And so this involved defining 36 classes (for each letter in the alphabet and digits 0–9.) and building a synthesized data set from EMNIST where multiple letters/digits are maintained and labeled on single images. Lastly, the tinyYOLOv3 darkflow framework was used to train the model to detect both the location and label of digits on over 400 images. Unfortunately, training the model was unsuccessful on *tinyYOLOv3*. However, worked on the pre-build library for *tiny – YOLOv4* made by Darkflow and Roboflow.

## 1 GPU - Hardware Requirements

Typically, YOLO frameworks including darkflow are run on dedicated NVIDIA GPU cards, this drastically increases training and query time. Here colab was used to remotely connect to a GPU, specifically, a Tesla P100-PCIE was used to train the model (Note, colab does not guarantee the quality of available GPU's and so training speed may vary).

---

```
1 gpu_info = !nvidia-smi
2 gpu_info = '\n'.join(gpu_info)
3 if gpu_info.find('failed') >= 0:
4     print('Not connected to a GPU')
5 else:
6     print(gpu_info)
```

---

## 2 Installing Darkflow And Tensor With Compatibility

Darkflow implementation of YOLO was retrained on the synthesized data set. For this, it is critical that all libraries packages are compatible with the version of darkflow/tinyYOLO.

---

```
1 !pip install tensorflow-gpu==1.15.0
2 !pip install imageio
3
4 # Download and build darkflow (the tensorflow implementation of YOLO)
5 import os
6 import pathlib
7
8 if "darkflow-master" in pathlib.Path.cwd().parts:
9     while "darkflow-master" in pathlib.Path.cwd().parts:
10         os.chdir('..')
11 elif not pathlib.Path("darkflow-master").exists():
12     !git clone --depth 1 https://github.com/thtrieu/darkflow.git
13     # Compile darkflow
14     %cd darkflow
15     !python setup.py build_ext --inplace
```

```

16 # Change darkflow to darkflow-master to distinguish between folder names
17 %cd ../
18 !mv darkflow darkflow-master
19 %cd darkflow-master
20
21
22 # Upload yolo.weights, pre-trained weights file (for YOLO v2) from an external Google drive
23 weights = 'yolo'
24 weights_file = weights + '.weights'
25 if not os.path.exists('weights_file'):
26     !gdown --id 0BtW_VtY7onTnBYYWdqSHNGSUU
27     !mkdir bin
28     !mv yolo.weights bin
29
30
31 # Imports
32 %cd darkflow-master
33 %tensorflow_version 1.15.0rc2
34
35 # For importing/exporting files, working with arrays, etc
36 import time
37 import urllib
38 import numpy as np
39 import pandas as pd
40 import imageio
41
42 # For actual object detection
43 import tensorflow as tf
44 from darkflow.net.build import TFNet
45 threshold = 0.25
46
47 # For drawing onto and plotting images
48 import matplotlib.pyplot as plt
49 import cv2
50 %config InlineBackend.figure_format = 'svg'

```

---

From here a model pipeline can be constructed by,

---

```

1 params = {
2     'model': 'cfg/yolo.cfg',
3     'load': 'bin/yolo.weights',
4     'threshold': threshold,
5     'gpu': 1.0
6 }
7
8 # Run the model
9 tfnet = TFNet(params)

```

---

### 3 Testing Darkflow Detection

The above implementation of darkflow has a pre-computed deep neural network for multiclass object detection which also automatically downloads demo images to validate if darkflow was armed correctly, (without incompatibility errors). And so this default test was run to check the model.

---

```

1 img = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
2 result = tfnet.return_predict(img)
3 for i in range(0, len(result)) :
4     img = cv2.rectangle(img, (result[i]['topleft']['x'], result[i]['topleft']['y']), (result[i]['bottomright']['x'], result[i]['bottomright']['y']), (0, 255, 0), 7)
5     img = cv2.putText(img, result[i]['label'], (result[i]['topleft']['x'], result[i]['topleft']['y']), cv2.FONT_HERSHEY_COMPLEX, 1, (0, 0, 0), 2)
6 plt.imshow(img)

```

---

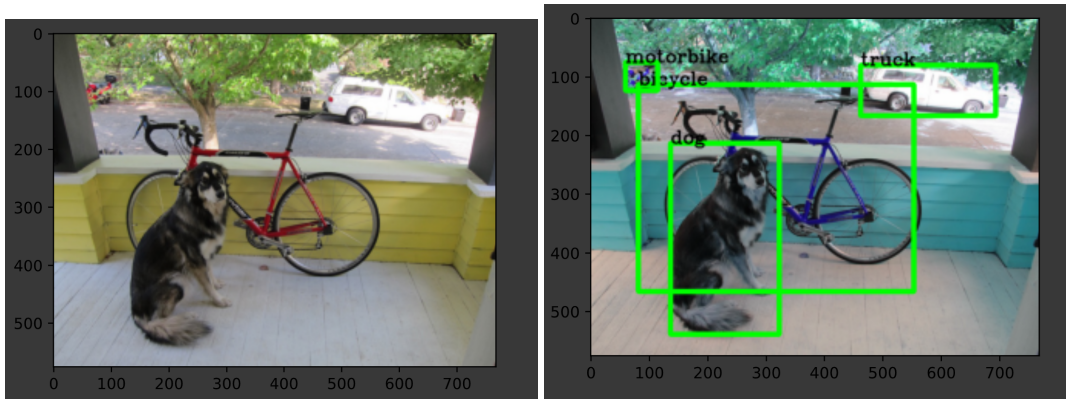


Figure 1: Darkflow - Master Demo of YOLOv3 Framework

## 4 Synthesizing Dataset EMNIST

EMNIST is a data set that contains  $28 \times 28$  images of both digits and letters. It is an extension created by, Gregory Cohen et. al. of the original MNIST which only had digits from 0 – 9. The original (non-synthesized data set) is implemented by pip install,

---

```

1 from sklearn import preprocessing
2 from sklearn.model_selection import train_test_split
3 from sklearn.datasets import fetch_openml
4
5 !pip install emnist
6 from emnist import list_datasets

```

---

Also the data set needs to be warmed up,

---

```

1 list_datasets()

```

---

Finally, the raw training and testing images and labels (annotations) can be parsed in.

---

```

1 from emnist import extract_training_samples
2 from emnist import extract_test_samples
3
4 test_1, test_labels_1 = extract_test_samples('letters')
5 train_1, train_labels_1 = extract_training_samples('letters')
6
7 train_labels_1 = list(np.asarray(train_labels_1) + 10)
8 test_labels_1 = list(np.asarray(test_labels_1) + 10)
9
10 test_2, test_labels_2 = extract_test_samples('digits')
11 train_2, train_labels_2 = extract_training_samples('digits')
12
13 Test_images, Test_labels = np.concatenate((test_1, test_2), axis=0), np.concatenate((test_labels_1,
14     test_labels_2), axis=0)
15 Train_images, Train_labels = np.concatenate((train_1, train_2), axis=0), np.concatenate((train_labels_1,
16     train_labels_2), axis=0)
17
18 print(Test_images.shape)
19 print(Test_labels.shape)
20
21 print(Train_images.shape)
22 print(Train_labels.shape)

```

---

Raw images in either set are typically bounded around the letter, Importantly, these images do not contain information about their bounding boxes and also only have a single image per image. And so are initially poor training candidates for the YOLO framework. As a result, we construct a new more compatible data set from the raw EMNIST images.

Here, a larger ( $200 \times 200$ ) black image is created, then randomly selected and randomly sized EMNIST data images are pasted/ overlaid on the blank image to make a larger image with multiple hand written digits. Critically, by implementing this data set in this way, the original labels from EMNIST are no long compatible (because they do not contain spacial information about where each letter is located in the larger image.) And so




Figure 2: EMNIST Data Image Examples

new annotations/labels are parse where each image contains a non trivial set of the original labels describing both the locations of the bounding boxes around each digit and their label, 1 – 36.




Figure 3: Synthesized Dataset Examples

Additionally, annotations needed to be in a compatible format for Darkflow automatic parsing. Which is .xml web based format. And so annotations were automatically generated to a corresponding file to each image. Each annotation contained information about each digit on the larger image, including bounding box dimensions and digit label.

---

```

1  import random
2  from PIL import Image
3  from xml.etree import ElementTree as ET
4
5  dataset_size = 1
6  for set_size in range(0,dataset_size) :
7
8      #Blank Image
9      temp_img = np.zeros([200,200],dtype=np.uint8)
10
11     index = set_size
12     annot = "<annotation>\n\
13             <folder>VOC2007</folder>\n\
14             <filename>" + str(index) + ".jpg</filename>\n\
15             <source>\n\
16             <database>EMNIST</database>\n\
17             <annotation>EMNIST</annotation>\n\
18             <image>EMNIST</image>\n\
19             </source>\n\
20             <owner>\n\
21             <name>Gregory Cohen, Saeed Afshar, Jonathan Tapson, and Andre van Schaik</name>\n\
22             </owner>\n\
23             <size>\n\
24             <width>200</width>\n\
25             <height>200</height>\n\
26             <depth>1</depth>\n\

```

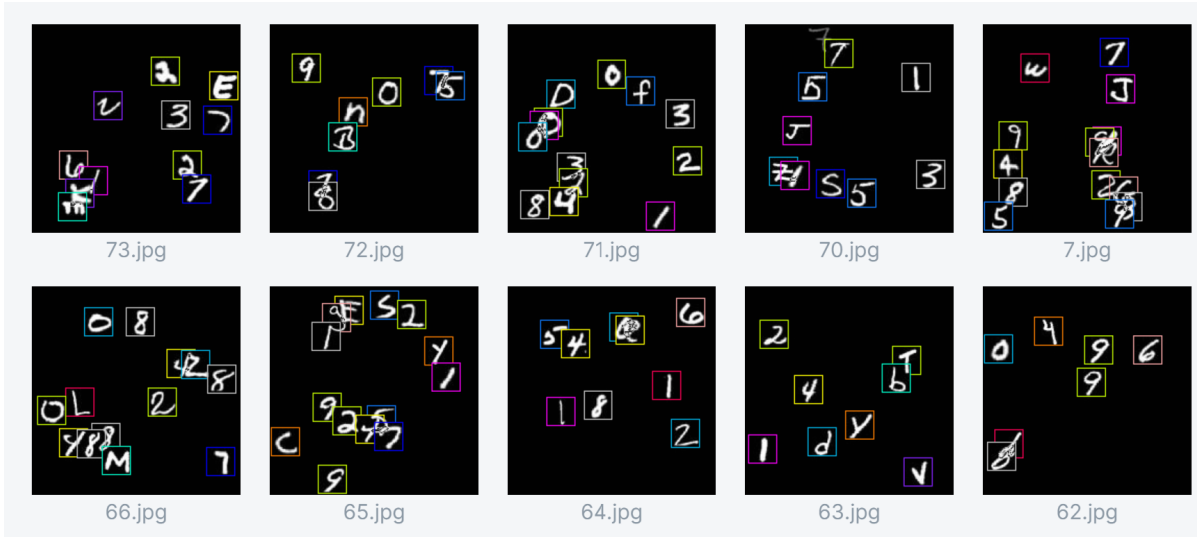


Figure 4: Synthesized Data set with bounding boxes

```

27         </size>\n\
28         <segmented>0</segmented>"
29     stack = random.randint(7, 15)
30     for i in range(0, stack):
31         j = random.randint(0, len(Train_images))
32         anchor = ( random.randint(0,200 - len(Train_images[j]-1)), random.randint(0,200-len(Train_images[j]
33             )-1))
34         temp_img[ anchor[0]:anchor[0]+len(Train_images[j]), anchor[1]:anchor[1]+len(Train_images[j]) ] +=
35             Train_images[j]
36         annot += "<object>\n\
37             <name>" + str(Train_labels[j]) + "</name>\n\
38             <pose>Left</pose>\n\
39             <truncated>0</truncated>\n\
40             <difficult>0</difficult>\n\
41             <bndbox>\n\
42                 <xmin>" + str(anchor[0]) + "</xmin>\n\
43                 <ymin>" + str(anchor[1]) + "</ymin>\n\
44                 <xmax>" + str(anchor[0] + len(Train_images[j])) + "</xmax>\n\
45                 <ymax>" + str(anchor[1] + len(Train_images[j])) + "</ymax>\n\
46             </bndbox>\n\
47             </object>\n\
48         annot += "</annotation>"
49
50     tree = ET.XML(annot)
51     with open("/content/darkflow-master/test/training/annotations/" + str(index) + '.xml', "wb") as f:
52         f.write(ET.tostring(tree))
53
54     im = Image.fromarray(temp_img)
55     im.save("/content/darkflow-master/test/training/images/" + str(index) + ".jpg")
56     #temp_img = temp_img.paste(img2, (0,0))

```

The dataset was generated with a subset of 40,000 randomly selected EMNIST images to create 200 larger images containing multiple digit objects. And including bounding boxes results in,

## 5 Training - Original Darkflow *tiny* – YOLOv3

After the data set was initialized, Darkflow's YOLO implementation can be re-trained for the new data, with the designated number of classes.

First, ensure the colab runtime is still connected to a GPU,

```

1 physical_devices = tf.config.experimental.list_physical_devices('GPU')
2 tf.config.experimental.set_memory_growth(physical_devices[0], True)

```

Then, train the model using the file locations that store the data-set

---

```
1 with tf.device('/device:GPU:0') :  
2     !python flow --model cfg/tiny-yolo-voc-3c.cfg --load bin/yolo.weights --train --annotation test/  
    training/annotations --dataset test/training/images
```

---

Unfortunately, this method was incompatible with colabs GPU and so the training time run on CPU exceeded my patience.

## 6 Training Roboflow Darkflow *tiny* – YOLOv4

An alternative training model was considered by uploading the synthesized dataset to Roboflow and the using the pre-built colab notebook by Darkflow that is directly compatible with Roboflow. This is mainly to test the validity of the data set and I did not write any of the code in the pre-built library. If anything was changed it was just the model parameters.

However, the code in the pre-build notebook follows the same logic as the above implementation.

Critically, this also involves switch darkflow versions to YOLOv4-tiny which is significantly faster than YOLOv3-tiny.

The model was trained for 30 minutes and these were some of the test results. **Note: that labels are numerical, 0 – 36, and so letters are labeled by increasing alphabet from 10. So *a* is 11, and *z* is 36.**





Figure 5: Test Output: tinyTOLOv4 model prediction



## 7 Appendix

---



---

```

1 gpu_info = !nvidia-smi
2 gpu_info = '\n'.join(gpu_info)
3 if gpu_info.find('failed') >= 0:
4     print('Not connected to a GPU')
5 else:
6     print(gpu_info)
7
8 from psutil import virtual_memory
9 ram_gb = virtual_memory().total / 1e9
10 print('Your runtime has {:.1f} gigabytes of available RAM\n'.format(ram_gb))
11
12 if ram_gb < 20:
13     print('Not using a high-RAM runtime')
14 else:
15     print('You are using a high-RAM runtime!')
16
17 Imports: Darkflow (YOLO) and Tensorflow
18
19 from sklearn import svm
20 import matplotlib.pyplot as plt
21 from sklearn import preprocessing
22 from sklearn.model_selection import train_test_split
23 from sklearn.datasets import fetch_openml
24 import numpy as np
25
26
27
28 from google.colab import drive
29 from pydrive.auth import GoogleAuth
30 from pydrive.drive import GoogleDrive
31 from google.colab import auth
32 from oauth2client.client import GoogleCredentials
33 from google.colab import drive
34 drive.mount('/content/drive')
35
36 auth.authenticate_user()
37 gauth = GoogleAuth()
38 gauth.credentials = GoogleCredentials.get_application_default()
39 drive = GoogleDrive(gauth)
40
41 # Install required libraries
42 !pip install tensorflow-gpu==1.15.0rc2
43 !pip install tensorflow-gpu==1.15.0
44 !pip install imageio

```

```

45
46 # Download and build darkflow (the tensorflow implementation of YOLO)
47 import os
48 import pathlib
49
50 if "darkflow-master" in pathlib.Path.cwd().parts:
51     while "darkflow-master" in pathlib.Path.cwd().parts:
52         os.chdir('..')
53 elif not pathlib.Path("darkflow-master").exists():
54     !git clone --depth 1 https://github.com/thtrieu/darkflow.git
55     # Compile darkflow
56     %cd darkflow
57     !python setup.py build_ext --inplace
58     # Change darkflow to darkflow-master to distinguish between folder names
59     %cd ../
60     !mv darkflow darkflow-master
61     %cd darkflow-master
62
63
64 # Upload yolo.weights, pre-trained weights file (for YOLO v2) from an external Google drive
65 weights = 'yolo'
66 weights_file = weights + '.weights'
67 if not os.path.exists('weights_file'):
68     !gdown --id 0BItW_VtY7onTnBYWdqSHNGSUU
69     !mkdir bin
70     !mv yolo.weights bin
71
72
73 # Imports
74 %cd darkflow-master
75 %tensorflow_version 1.15.0rc2
76
77 # For importing/exporting files, working with arrays, etc
78 import time
79 import urllib
80 import numpy as np
81 import pandas as pd
82 import imageio
83
84 # For actual object detection
85 import tensorflow as tf
86 from darkflow.net.build import TFNet
87 threshold = 0.25
88
89 # For drawing onto and plotting images
90 import matplotlib.pyplot as plt
91 import cv2
92 %config InlineBackend.figure_format = 'svg'
93
94 Define a Model:
95
96 params = {
97     'model': 'cfg/yolo.cfg',
98     'load': 'bin/yolo.weights',
99     'threshold': threshold,
100     'gpu': 0.0
101 }
102
103 # Run the model
104 tfnet = TFNet(params)
105
106 Test a image and make a bounding box
107
108 import matplotlib.pyplot as plt
109 import matplotlib.image as mp
110 import PIL as pl
111 import cv2
112
113 image_path = '/content/darkflow-master/sample_img/sample_dog.jpg'
114 image = mp.imread(image_path)
115 plt.imshow(image)
116
117 Query Image

```



```

118
119 img = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
120 result = tfnet.return_predict(img)
121 img.shape
122
123
124 Extract Data
125
126 tl = (result[1][ 'topleft' ][ 'x' ], result[1][ 'topleft' ][ 'y' ])
127 br = (result[1][ 'bottomright' ][ 'x' ], result[1][ 'bottomright' ][ 'y' ])
128 label = result[1][ 'label' ]
129
130 print(label)
131
132 Display Box
133
134 for i in range(0, len(result)) :
135     img = cv2.rectangle(img, (result[i][ 'topleft' ][ 'x' ], result[i][ 'topleft' ][ 'y' ]), (result[i][ 'bottomright' ][ 'x' ],
136         result[i][ 'bottomright' ][ 'y' ]), (0, 255, 0), 7)
137     img = cv2.putText(img, result[i][ 'label' ], (result[i][ 'topleft' ][ 'x' ], result[i][ 'topleft' ][ 'y' ]), cv2.
138         FONT_HERSHEY_COMPLEX, 1, (0, 0, 0), 2)
139 plt.imshow(img)
140
141 # EMNIST Data
142
143 from sklearn import preprocessing
144 from sklearn.model_selection import train_test_split
145 from sklearn.datasets import fetch_openml
146
147 !pip install emnist
148
149 from emnist import list_datasets
150
151 list_datasets()
152
153 from emnist import extract_training_samples
154 from emnist import extract_test_samples
155
156 test_1, test_labels_1 = extract_test_samples('letters')
157 train_1, train_labels_1 = extract_training_samples('letters')
158
159 train_labels_1 = list(np.asarray(train_labels_1) + 10)
160 test_labels_1 = list(np.asarray(test_labels_1) + 10)
161
162 test_2, test_labels_2 = extract_test_samples('digits')
163 train_2, train_labels_2 = extract_training_samples('digits')
164
165 Test_images, Test_labels = np.concatenate((test_1, test_2), axis=0), np.concatenate((test_labels_1, test_labels_2),
166     axis=0)
167 Train_images, Train_labels = np.concatenate((train_1, train_2), axis=0), np.concatenate((train_labels_1,
168     train_labels_2), axis=0)
169
170 print(Test_images.shape)
171 print(Test_labels.shape)
172
173 print(Train_images.shape)
174 print(Train_labels.shape)
175
176 Train_images.shape
177
178 train_1.shape
179
180 train_2.shape
181
182 Train_images[55].shape
183
184 #Let's print first 10 letters from EMNIST training samples
185 for i in range(7):
186     first_ten_letters = Train_images[i].reshape((28, 28))
187     plt.imshow(first_ten_letters, cmap = 'binary')
188     #plt.axis("off")
189     plt.show()

```

```

187
188
189 # Gonna Synthesize my out data set with boxes...
190
191 plt.imshow(Train_images[5])
192
193 print(Train_labels[:6])
194
195 print(len(Train_images[50]))
196
197 print(str(10))
198
199
200 import random
201 from PIL import Image
202 from xml.etree import ElementTree as ET
203
204 dataset_size = 80
205 for set_size in range(0, dataset_size) :
206
207     #Blank Image
208     temp_img = np.zeros([200,200], dtype=np.uint8)
209
210     index = set_size
211     annot = "<annotation>\n\
212         <folder>VOC2007</folder>\n\
213         <filename>" + str(index) + ".jpg</filename>\n\
214         <source>\n\
215         <database>EMNIST</database>\n\
216         <annotation>EMNIST</annotation>\n\
217         <image>EMNIST</image>\n\
218         </source>\n\
219         <owner>\n\
220         <name>Gregory Cohen, Saeed Afshar, Jonathan Tapson, and Andre van Schaik</name>\n\
221         </owner>\n\
222         <size>\n\
223         <width>200</width>\n\
224         <height>200</height>\n\
225         <depth>1</depth>\n\
226         </size>\n\
227         <segmented>0</segmented>"
228     stack = random.randint(7, 15)
229     for i in range(0, stack) :
230         j = random.randint(0, len(Train_images))
231         anchor = ( random.randint(0,200 - len(Train_images[j])-1), random.randint(0,200-len(Train_images[j])-1))
232         temp_img[anchor[0]:anchor[0]+len(Train_images[j]), anchor[1]:anchor[1]+len(Train_images[j])] += Train_images[j]
233         annot += "<object>\n\
234             <name>" + str(Train_labels[j]) + "</name>\n\
235             <pose>Left</pose>\n\
236             <truncated>0</truncated>\n\
237             <difficult>0</difficult>\n\
238             <bndbox>\n\
239             <xmin>" + str(anchor[1]) + "</xmin>\n\
240             <ymin>" + str(anchor[0]) + "</ymin>\n\
241             <xmax>" + str(anchor[1] + len(Train_images[j])) + "</xmax>\n\
242             <ymax>" + str(anchor[0] + len(Train_images[j])) + "</ymax>\n\
243             </bndbox>\n\
244             </object>\n\
245         </annotation>"
246
247     tree = ET.XML(annot)
248     with open("/content/darkflow-master/test/training/annotations/" + str(index) + '.xml', "wb") as f :
249         f.write(ET.tostring(tree))
250
251     im = Image.fromarray(temp_img)
252     im.save("/content/darkflow-master/test/training/images/" + str(index) + ".jpg")
253     #temp_img = temp_img.paste(img2, (0,0))
254
255 # Training
256
257 physical_devices = tf.config.experimental.list_physical_devices('GPU')
258 tf.config.experimental.set_memory_growth(physical_devices[0], True)
259

```

```

260 print("Nm GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
261
262 tf.test.gpu_device_name()
263
264 tf.config.list_physical_devices('GPU')
265
266 with tf.device('/device:GPU:0') :
267     !python flow --model cfg/tiny-yolo-voc-3c.cfg --load bin/yolo.weights --train --annotation test/training/
        annotations --dataset test/training/images
268
269 !zip -r /content/darkflow-master/test/training.zip /content/darkflow-master/test/training
270
271 from google.colab import files
272 files.download('/content/darkflow-master/test/training.zip')
273
274
275
276 files.download("/content/darkflow-master/test/training")

```

---